ACL 2020 - Select Papers List

DISCLAIMER: This is a select list of papers presented at ACL - 2020 conference. Not intended to be a representative list for the conference or any organization

ACL 2020 Papers:

Conference Proceedings: https://www.aclweb.org/anthology/events/acl-2020

Best Overall Paper: Beyond Accuracy: Behavioral Testing of NLP Models with CheckList

Honorable Mentions:

- Tangled up in BLEU: Reevaluating the Evaluation of Automatic Machine Translation Evaluation Metrics
- Don't Stop Pretraining: Adapt Language Models to Domains and Tasks

Machine Learning for NLP:

- Weight Poisoning Attacks on Pretrained Models
- Do Transformers Need Deep Long-Range Memory?
- Pretrained Transformers Improve Out-of-Distribution Robustness
- MobileBERT: a Compact Task-Agnostic BERT for Resource-Limited Devices
- Low Resource Sequence Tagging using Sentence Reconstruction
- Zero-shot Text Classification via Reinforced Self-training
- Why Overfitting Isn't Always Bad: Retrofitting Cross-Lingual Word Embeddings to Dictionaries

Interpretability and Analysis of Models for NLP:

- On the Cross-lingual Transferability of Monolingual Representations
- Finding Universal Grammatical Relations in Multilingual BERT
- Towards Transparent and Explainable Attention Models
- Human Attention Maps for Text Classification: Do Humans and Neural Networks Focus on the Same Words?
- exBERT: A Visual Analysis Tool to Explore Learned Representations in Transformer Models

Ethics in NLP:

- Gender Bias in Multilingual Embeddings and Cross-Lingual Transfer
- It's Morphin' Time! Combating Linguistic Discrimination with Inflectional Perturbations
- Double-Hard Debias: Tailoring Word Embeddings for Gender Bias Mitigation
- Towards Debiasing Sentence Representations

NLP Applications:

- DeeBERT: Dynamic Early Exiting for Accelerating BERT Inference
- Improving Segmentation for Technical Support Problems

Information Extraction:

- Named Entity Recognition without Labelled Data: A Weak Supervision Approach
- Improving Low-Resource Named Entity Recognition using Joint Sentence and Token Labeling
- Single-/Multi-Source Cross-Lingual NER via Teacher-Student Learning on Unlabeled Data in Target Language
- NAT: Noise-Aware Training for Robust Neural Sequence Labeling

Dialogue and Interactive Systems:

- Gated Convolutional Bidirectional Attention-based Model for Off-topic Spoken Response Detection
- Negative Training for Neural Dialogue Response Generation
- A Generative Model for Joint Natural Language Understanding and Generation
- Learning Low-Resource End-To-End Goal-Oriented Dialog for Fast and Reliable System Deployment

Machine Translation:

- BPE-Dropout: Simple and Effective Subword Regularization
- Unsupervised Word Translation with Adversarial Autoencoder
- Selecting Backtranslated Data from Multiple Sources for Improved Neural Machine Translation
- Multi-Domain Neural Machine Translation with Word-Level Adaptive Layer-wise Domain Mixing

Question Answering:

- Injecting Numerical Reasoning Skills into Language Models
- Selective Question Answering under Domain Shift

Speech and Multimodality:

- Learning Spoken Language Representations with Neural Lattice Language Modeling
- Towards Emotion-aided Multi-modal Dialogue Act Classification

Summarization:

- Attend, Translate and Summarize: An Efficient Method for Neural Cross-Lingual Summarization
- Improving Truthfulness of Headline Generation

Semantics:

- Curriculum Learning for Natural Language Understanding
- Mind the Trade-off: Debiasing NLU Models without Degrading the In-distribution Performance
- Towards Robustifying NLI Models Against Lexical Dataset Biases
- End-to-End Bias Mitigation by Modelling Biases in Corpora
- Syntactic Data Augmentation Increases Robustness to Inference Heuristics

Computational Social Science and Social Media:

- Text and Causal Inference: A Review of Using Text to Remove Confounding from Causal Estimates